

Problem Set 2 {Supervised Learning II}

1. Logistic Regression: Training Stability

Goal: Develop your skills debugging machine Learning Algorithms.

→ Implementation of Logistic Regression is provided in `src/pol_lr.py`.

→ Two labeled dataset:

① `data/ds1-a.txt`

② `data/ds1-b.txt`

① * On dataset A, the training converges in 2,78,103 iterations.

* On dataset B, the training did not converge until 50,00,00 iterations

↳ Maybe it never converges

② $x, red \rightarrow y = 0$
 $o, blue \rightarrow y = 1$

learning_rate = 1

$\|\Delta\theta\| < 10^{-15}$ for convergence

Observation

1) for dataset A ~~at algorithm~~ in the learning phase, θ tends to converge to a fixed point.

2) for dataset B, in the learning phase, θ tends to continuously increase.

⇒ If the dataset can be perfectly separated by a linear decision boundary, then the objective $J(\theta)$ can be arbitrarily ~~maximized~~ minimized by just scaling θ .

→ Multiplying θ by a factor does not change the decision boundary.

→ But probability being assigned to each data point can be arbitrarily changed

$$h_{\theta}(A) = \frac{1}{1 + e^{-\theta^T A}}$$

If $\theta \rightarrow k\theta$

$$h_{\theta}^*(A) = \frac{1}{1 + e^{-k\theta^T A}}$$

$$h_{\theta}(A) < h_{\theta}^*(A)$$


$$\Rightarrow \text{as } k \rightarrow \infty \quad h_{\theta}^*(A) \rightarrow 1$$

ds1-a.txt

→ Cannot be separated perfectly, so we are able to converge at some θ .

ds2-b.txt

→ Can be separated perfectly, so we are not able to converge at any θ .



©

(i) No, there is no maxima in $l(\theta)$

~~(ii) No, there is no maxima in $l(\theta)$~~

~~(ii)~~ (ii) Yes, When the learning rate is sufficiently small, the update to θ will be small, and it will be judge to be converged by the algorithm.

(iii) No, The dataset will still be linearly separable.

(iv) Yes, In this case there is a cost of increasing θ .

(v) Yes, this will very likely make the dataset not linearly separable.

① No, SVM's objective is directly associated to geometric margin, which is independent with the scaling of θ .